

tf-quick-intro

April 19, 2020

1 A Quick Introduction to TensorFlow 2.0 for Deep Learning

1.1 <https://towardsdatascience.com/a-quick-introduction-to-tensorflow-2-0-for-deep-learning-e740ca2e974c>

1.2 by George Seif, Oct 24, 2019

1.2.1 Google Colab Notebook can be found at: <https://colab.research.google.com/drive/1KRi0k5X>

1.2.2 YP 200330

```
[2]: import tensorflow as tf
from tensorflow.keras import datasets, layers, models, optimizers
import matplotlib.pyplot as plt

if tf.__version__ < "2.0.0":
    !pip install --upgrade tensorflow_gpu==2.0
```

1.3 Loading data and visualisation

1.3.1 MNIST dataset contains


- * 60,000 training images
- * 10,000 test images
- * digits from 0 to 9
- * size 28x28

```
[3]: (train_images, train_labels), (test_images, test_labels) = datasets.mnist.
    ↪load_data()

class_names = ['zero', 'one', 'two', 'three', 'four', 'five',
               'six', 'seven', 'eight', 'nine']

plt.figure(figsize=(10,10))
for i in range(16):
    plt.subplot(4, 4, i+1)
```

```
plt.xticks([])
plt.yticks([])
plt.imshow(train_images[i], cmap=plt.cm.binary)
plt.xlabel(class_names[train_labels[i]])
plt.show()
```



1.4 Model Definition

```
[4]: IMG_SIZE = (28, 28, 1)
input_img = layers.Input(shape=IMG_SIZE)

model = layers.Conv2D(32, (3, 3), padding='same')(input_img)
```

```

model = layers.Activation('relu')(model)
model = layers.Conv2D(32, (3, 3), padding='same', strides=(2, 2))(model)
model = layers.Activation('relu')(model)

model = layers.Conv2D(64, (3, 3), padding='same')(model)
model = layers.Activation('relu')(model)
model = layers.Conv2D(64, (3, 3), padding='same', strides=(2, 2))(model)
model = layers.Activation('relu')(model)

model = layers.Conv2D(64, (3, 3), padding='same')(model)
model = layers.Activation('relu')(model)
model = layers.Conv2D(64, (3, 3), padding='same')(model)
model = layers.Activation('relu')(model)

model = layers.GlobalAveragePooling2D()(model)
model = layers.Dense(32)(model)
model = layers.Activation('relu')(model)
model = layers.Dense(10)(model)
output_img = layers.Activation('softmax')(model)

model = models.Model(input_img, output_img)

model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d (Conv2D)	(None, 28, 28, 32)	320
activation (Activation)	(None, 28, 28, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 32)	9248
activation_1 (Activation)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 64)	18496
activation_2 (Activation)	(None, 14, 14, 64)	0
conv2d_3 (Conv2D)	(None, 7, 7, 64)	36928
activation_3 (Activation)	(None, 7, 7, 64)	0
conv2d_4 (Conv2D)	(None, 7, 7, 64)	36928

```

-----
activation_4 (Activation)      (None, 7, 7, 64)           0
-----
conv2d_5 (Conv2D)             (None, 7, 7, 64)          36928
-----
activation_5 (Activation)      (None, 7, 7, 64)           0
-----
global_average_pooling2d (Gl (None, 64)           0
-----
dense (Dense)                 (None, 32)                 2080
-----
activation_6 (Activation)      (None, 32)                 0
-----
dense_1 (Dense)               (None, 10)                 330
-----
activation_7 (Activation)      (None, 10)                 0
=====
Total params: 141,258
Trainable params: 141,258
Non-trainable params: 0
-----

```

1.5 Training and Testing

```

[5]: train_images = train_images.reshape(60000, 28, 28, 1).astype('float32') / 255.0
      test_images = test_images.reshape(10000, 28, 28, 1).astype('float32') / 255.0

      train_labels = tf.keras.utils.to_categorical(train_labels, 10)
      test_labels = tf.keras.utils.to_categorical(test_labels, 10)

```

```

[6]: adam = optimizers.Adam(lr=0.0001)
      model.compile(adam, loss='categorical_crossentropy', metrics=["accuracy"])

      history = model.fit(train_images, train_labels, epochs=10,
                          validation_data=(test_images, test_labels))

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 7s 124us/sample - loss: 0.7584 -
accuracy: 0.7474 - val_loss: 0.3452 - val_accuracy: 0.9013

Epoch 2/10

60000/60000 [=====] - 6s 103us/sample - loss: 0.2451 -
accuracy: 0.9297 - val_loss: 0.1927 - val_accuracy: 0.9433

Epoch 3/10

60000/60000 [=====] - 6s 106us/sample - loss: 0.1698 -
accuracy: 0.9506 - val_loss: 0.1262 - val_accuracy: 0.9647

Epoch 4/10

```

60000/60000 [=====] - 6s 104us/sample - loss: 0.1330 -
accuracy: 0.9606 - val_loss: 0.0880 - val_accuracy: 0.9726
Epoch 5/10
60000/60000 [=====] - 6s 105us/sample - loss: 0.1091 -
accuracy: 0.9674 - val_loss: 0.0942 - val_accuracy: 0.9714
Epoch 6/10
60000/60000 [=====] - 7s 109us/sample - loss: 0.0923 -
accuracy: 0.9729 - val_loss: 0.0848 - val_accuracy: 0.9734
Epoch 7/10
60000/60000 [=====] - 6s 108us/sample - loss: 0.0824 -
accuracy: 0.9747 - val_loss: 0.0760 - val_accuracy: 0.9750
Epoch 8/10
60000/60000 [=====] - 6s 105us/sample - loss: 0.0715 -
accuracy: 0.9786 - val_loss: 0.0763 - val_accuracy: 0.9750
Epoch 9/10
60000/60000 [=====] - 6s 105us/sample - loss: 0.0637 -
accuracy: 0.9805 - val_loss: 0.0503 - val_accuracy: 0.9835
Epoch 10/10
60000/60000 [=====] - 6s 105us/sample - loss: 0.0564 -
accuracy: 0.9831 - val_loss: 0.0584 - val_accuracy: 0.9817

```

```


[7]: plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.7, 1])
plt.legend(loc='best')

```

```


[7]: <matplotlib.legend.Legend at 0x7f77504c8e90>

```



```
[8]: plt.plot(history.history['loss'], label='loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(loc='best')
```

```
[8]: <matplotlib.legend.Legend at 0x7f775049f110>
```



```
[9]: test_loss, test_accuracy = model.evaluate(test_images, test_labels, verbose=2)
print('\nTest accuracy = {0:.2f}%'.format(test_accuracy*100.0))
```

10000/1 - 0s - loss: 0.0293 - accuracy: 0.9817

Test accuracy = 98.17%